

# The iVM assembler and prototype machine implementations (architecture version 2.0, may 2023)

Source: <https://github.com/immortalvm/ivm-implementations>

## Introduction

The iVM assembler was made for two reasons: To make it feasible to program the iVM by hand, and also to act as a target for the iVM compiler. The assembly language provides some useful shorthands and it abstracts away details that does not regard the assembly programmer or the compiler.

## The assembly language: semi-formal EBNF, ignoring whitespace and comments

```
program = import* statement* ;
```

```
import = "IMPORT" (identifier "/" )+ identifier
```

```
statement = identifier ":" (* label *)
  | "EXPORT" identifier (* export declaration *)
  | identifier "=" expression (* abbreviation *)
  | "data1" "[" expression* "]" ("*" positive_numeral)? (* data segment, 8 bits per value *)
  | "data2" "[" expression* "]" ("*" positive_numeral)? (* data segment, 16 bits per value *)
  | "data4" "[" expression* "]" ("*" positive_numeral)? (* data segment, 32 bits per value *)
  | "data8" "[" expression* "]" ("*" positive_numeral)? (* data segment, 64 bits per value *)
  | "space" expression (* pointer static byte array *)

  | "exit" | "exit!" expression
  | "push" | "push!" expression | "push!!" expression expression | ...
  | "set_sp" | "set_sp!" expression
  | "jump" | "jump!" expression
  | "jump_zero" | "jump_zero!" expression | "jump_zero!!" expression expression
  | "jump_not_zero" | "jump_not_zero!" expression | "jump_not_zero!!" expression expression
  | "call" | "call!" expression
  | "return" (* alias for "jump" *)
  | "check_version"

  | "load1" | "load1!" expression
  | "load2" | "load2!" expression
  | "load4" | "load4!" expression
  | "load8" | "load8!" expression
  | "sigx1" | "sigx1!" expression
  | "sigx2" | "sigx2!" expression
  | "sigx4" | "sigx4!" expression
  | "sigx8" | "sigx8!" expression (* no-op *)
  | "store1" | "store1!" expression | "store1!!" expression expression
  | "store2" | "store2!" expression | "store2!!" expression expression
  | "store4" | "store4!" expression | "store4!!" expression expression
  | "store8" | "store8!" expression | "store8!!" expression expression

  | "add" | "add!" expression | "add!!" expression expression
  | "sub" | "sub!" expression | "sub!!" expression expression
  | "mult" | "mult!" expression | "mult!!" expression expression
  | "neg" | "neg!" expression
  | "and" | "and!" expression | "and!!" expression expression
  | "or" | "or!" expression | "or!!" expression expression
  | "xor" | "xor!" expression | "xor!!" expression expression
  | "not" | "not!" expression
  | "pow2" | "pow2!" expression
```

```
| "shi ft_l" | "shi ft_l!" expressi on | "shi ft_l!!" expressi on expressi on
| "shi ft_ru" | "shi ft_ru!" expressi on | "shi ft_ru!!" expressi on expressi on
| "shi ft_rs" | "shi ft_rs!" expressi on | "shi ft_rs!!" expressi on expressi on
```

```
| "di v_u" | "di v_u!" expressi on | "di v_u!!" expressi on expressi on
| "di v_s" | "di v_s!" expressi on | "di v_s!!" expressi on expressi on
| "rem_u" | "rem_u!" expressi on | "rem_u!!" expressi on expressi on
| "rem_s" | "rem_s!" expressi on | "rem_s!!" expressi on expressi on
```

```
| "l t_u" | "l t_u!" expressi on | "l t_u!!" expressi on expressi on
| "l t_s" | "l t_s!" expressi on | "l t_s!!" expressi on expressi on
| "l te_u" | "l te_u!" expressi on | "l te_u!!" expressi on expressi on
| "l te_s" | "l te_s!" expressi on | "l te_s!!" expressi on expressi on
| "eq" | "eq!" expressi on | "eq!!" expressi on expressi on
| "gte_u" | "gte_u!" expressi on | "gte_u!!" expressi on expressi on
| "gte_s" | "gte_s!" expressi on | "gte_s!!" expressi on expressi on
| "gt_u" | "gt_u!" expressi on | "gt_u!!" expressi on expressi on
| "gt_s" | "gt_s!" expressi on | "gt_s!!" expressi on expressi on
```

```
| "read_char" (* interactive *)
| "read_frame" | "read_frame!" expressi on
| ... | "read_pi xel!!" expressi on expressi on
```

```
| ... | "put_char!" expressi on
| ... | "put_byte!" expressi on
| ... | "new_frame!!!" expressi on expressi on expressi on
| ... | "set_pi xel!!!!" expressi on ...
| ... | "add_sample!!" expressi on expressi on;
```

```
expressi on = posi tive_numeral (* 0 to 2^64-1 *)
| identi fier (* label or abbrevi ation *)
```

```
| "-" expressi on (* correspondi ng statement: neg *)
| "~" expressi on (* not *)
| "$" expressi on (* stack content *)
| "&" expressi on (* stack pointer *)
```

```
| "(" "+" expressi on* ")" (* add *)
| "(" "*" expressi on* ")" (* mul t *)
| "(" "&" expressi on* ")" (* and *)
| "(" "|" expressi on* ")" (* or *)
| "(" "^" expressi on* ")" (* xor *)
```

```
| "(" "=" expressi on expressi on ")" (* eq *)
| "(" "<u" expressi on expressi on ")" (* l t_u *)
| "(" "<s" expressi on expressi on ")" (* l t_s *)
| "(" "<=u" expressi on expressi on ")" (* l te_u *)
| "(" "<=s" expressi on expressi on ")" (* l te_s *)
| "(" ">u" expressi on expressi on ")" (* gt_u *)
| "(" ">s" expressi on expressi on ")" (* gt_s *)
| "(" ">=u" expressi on expressi on ")" (* gte_u *)
| "(" ">=s" expressi on expressi on ")" (* gte_s *)
```

```
| "(" "<<" expressi on expressi on ")" (* shi ft_l *)
| "(" ">>u" expressi on expressi on ")" (* shi ft_r unsigned, unsigned *)
| "(" ">>s" expressi on expressi on ")" (* shi ft_r signed, signed *)
| "(" "/u" expressi on expressi on ")" (* di v_u *)
| "(" "/s" expressi on expressi on ")" (* di v_s *)
```

```

| "(" "%u" expression expression ")" (* rem_u *)
| "(" "%s" expression expression ")" (* rem_s *)

| "(" "load1" expression ")"
| "(" "load2" expression ")"
| "(" "load4" expression ")"
| "(" "load8" expression ")"
| "(" "sigx1" expression ")"
| "(" "sigx2" expression ")"
| "(" "sigx4" expression ")"
| "(" "sigx8" expression ")" (* identity function *)

```

```

identifier = (letter | "_" | "." ) (letter | "_" | "." | digit)*;

```

In v0.8 we added an alternative notation for “immediate arguments”:

```

push* [ <e1> <e2> ... <en> ]

```

is syntactic sugar for:

```

push!!! <e1> <e2> ... <en>          # with n exclamation marks

```

Similarly for the other statements, e.g. set\_pixel\*.

The arguments to space and data<N> must be compile time constants, except that data8 also accepts labels.

## Examples

### Simple assembly

```

### iVM assembly language introduction.
###
### Part 1 - Statements
###
### This file explains the iVM assembly language. It is itself a valid assembly
### file, but the code does not make much sense. Notice that # indicates that
### the rest of the line is a comment.
###
### Other than comments, an iVM assembly file consist of a list of statements.
### Whitespace is not significant, but it recommended to put each statement on a
### separate line.
###
### The assembly language is case-sensitive, with all the instructions written
### in lower case (even though we have used upper case in the headings for them
### to stand out).

### 1. SPECIAL STATEMENTS

## There are six special statements: imports, labels, exports, definitions,
## data and space statements.

## Import statements can only occur at the top of the file. It means that a
## label in file can be referenced below (provided that the label was
## exported in the other file). Circular dependencies are not allowed.
IMPORT intro2_basics/x

## A label statement indicates a place in the code (memory address) at
## runtime. By convention, all other statements should be indented.
my_label:

## For a label to be visible from other files, it must be "exported".
EXPORT my_label

## Definitions define abbreviations, usually constants.
prime_number = 982451653

## A label can be exported under a different name as follows:
external_name = my_label

```

```

EXPORT external_name

## Labels and definitions use the same namespace, which is independent from
## the names of instructions. Thus, you are free to define a label called
## add or exit. Names of labels and definitions consist of letters,
## digits and underscore (_), and they cannot start with a digit.

## Data statements specify bytes that should be included in the binary as is
## (more or less). It includes a whitespace-separated list of constant
## expressions, usually numbers. All numbers can be specified in decimal,
## octal or hexadecimal notation.

## data1 includes a list of bytes in the binary. Only the 8 least
## significant bits are used of each number. Thus, the last number can be
## replaced by 1.
data1 [ 0 1 -2 0o200 -0x99ff ]

## Similarly, data2 includes a list of (little-endian) 16-bit words, using
## the 16 least significant bits.
data2 [ 0x1000 0x2000 0x3000 ]

## data4 and data8 include lists of 32-bit and 64-bit words, respectively.
data4 [ 0x40000000 ]
data8 [-0x0123456789abcdef]

## It should be noted that starting a program with a data block is generally
## a bad idea, as our VM executes programs from the top. (Incidentally, 0
## means that the VM should terminate immediately. Thus, the meaningless
## statements below do not cause the machine to crash.)

## A space statement inserts a 64-bit pointer to memory allocated by the
## program at startup. The argument specifies the number of bytes that will
## be allocated.
my_1000_byte_array:
  space 1000

## The remaining statements correspond to actual machine instructions.
## However, there is not a one-to-one correspondence. The assembly language
## also contain several "pseudo-instructions" that translate into multiple
## native instructions. Moreover, the assembler will handle some technical
## issues such as choosing between long and short conditional jumps.

### 2. PUSH

## The push statement pushes 64-bit numbers onto the stack.
push! 13 # Push the number 13 onto the stack.

## As above, these numbers can be in either decimal, octal or hexadecimal
## notation; and they can be both positive and negative, with wrapping.
push! -1 # Push 0xffffffffffffff onto the stack.

## A single push statement can push multiple numbers, indicated by the
## number of exclamation marks (!).
push!! 0 1 # Push 0 onto the stack, then 1.
push # Do nothing

## The push statement is not only used to push constants.
push! my_label # Push the address of my_label.
push! prime_number # Push the result of expanding its definition.

n = 7 # An arbitrary number
push! &n # Push the address PC + n * 8 (of the nth element on the stack).
push! $n # Push the (64-bit) value at PC + n * 8.

## It is also possible to push the value of complex (Lisp style)
## expressions.
push! (+ my_label -$0) # Push the label address minus the top stack element.

## This will be explained in detail later, but observe that
#
## push!! 17 $0 is equivalent to push! 17 push! $1
#

```

```

## since within a statement $x and &x will be relative to value of the stack
## pointer at the start of the statement.

## The following notation is more convenient when pushing many values onto
## the stack:
push* [1 2 3 4 5 6 7]

## This is syntactic sugar for:
push!!!!!! 1 2 3 4 5 6 7

### 3. JUMP, JUMP_ZERO, JUMP_NOT_ZERO

## The jump statement changes the program counter to an address popped from
## the stack.
jump

## jump! x is (syntactic) sugar for push! x jump .
jump! my_label # Jump to my_label

## The jump_zero statement pops an address first, then a value, and jumps to
## the address if the value is zero.
jump_zero

## jump_zero! x is sugar for push! x jump_zero .
jump_zero! my_label # Jump to my_label if value popped is zero.

## Similarly, jump_zero!! x y is sugar for push!! x y jump_zero .
## The following statement jumps to my_label if the fourth element on the
## stack (counting from 0) is equal to prime_number.
jump_zero!! (+ prime_number -4) my_label

## The jump_not_zero statement is similar.
jump_not_zero

### 4. CALL, RETURN

## When calling a subroutine we want to continue execution at the next statement.
call! my_label # Sugar for push! <fresh> jump! my_label <fresh>:
return # Alias for jump

### 5. LOAD, SIGX

## The four load statements all pop one address from the stack and push a 64
## bit value, the contents of the memory at that location (and onwards).
## When less than 8 bytes are fetched, the value is considered unsigned and
## is 0-padded before pushing. We use little-endian encoding of multi-byte
## values.
load1 # Pop address and push the unsigned byte at that memory location.
load2 # ... unsigned 16-bit word ...
load4 # ... unsigned 32-bit word ...
load8 # Pop address and push the 64-bit word at that memory location.

## load! x is sugar for push! x load! , similarly for load2/4/8.
load4! my_label # Load the unsigned 32-bit word at my_label.

## If you want the signed value at a memory location instead, follow the load
## statement with a corresponding sigx statement (sigx1, sigx2 or sigx4).
#
## sigx1 performs "sign extension" from 8 to 64 bits. More precisely, it
## (1) pops a 64-bit value from the stack,
## (2) sets bits 8..63 (counting from 0) to the same state as bit 7,
## (3) pushes the result back onto the stack.
#
## sigx2 and sigx4 are similar.
sigx4 # Sign extension from 32 to 64 bits.

## For completeness, we also include the no-op sigx8, and let sigxN! x be
## sugar for push! x sigxN .
sigx8
sigx1! 0xff # Push -1.

```

### ### 6. STORE

```
## The four store operations pop an address A, then a value V from the stack
## and writes the least significant bytes of V to A.
store1      # Write the least significant byte of V to A.
store2      # Write 16 bits of V to A.
store4      # Write 32 bits of V to A.
store8      # Write all of V to A.

## storeN! x is sugar for push! x storeN
store4! my_label # Pop value and write lower 32 bits to memory at my_label.

## storeN!! x y is sugar for push!! x y storeN
store8!! prime_number my_label # Write prime_number to memory at my_label.
```

### ### 7. ARITHMETIC OPERATIONS

```
xx = 99
yy = -13

add          # Pop two values and push their sum (with wrapping).

## All arithmetic statements have corresponding "sugared" variants.
add! xx      # Sugar for push! xx add .
add!! xx yy  # Sugar for push!! xx yy add .

sub          # Pop x, then y, and push y - x. Notice the order!
sub! xx      # Subtract xx from the value on top of the stack.
sub!! xx yy  # Push xx - yy.

mult         # Pop two values and push their product.
neg          # Pop value and push its additive inverse (two's complement).

div_u        # Pop x, then y, and push y / x using unsigned division.
div_s        # Similar, but using signed division.

rem_u        # Pop two values and push the remainder of unsigned division.
rem_s        # Similar, but using signed division.
```

### ### 8. BITWISE OPERATIONS

```
and          # Pop two (64-bit) values and push their binary "and".

## All bitwise statements have corresponding "sugared" variants.
and! 0x7f    # Sugar for push! 0x7f and .
and!! 0xffff prime_number # Sugar for push!! 0xffff prime_number .

or           # Pop two values and push their binary "or".
xor          # Pop two values and push their binary "xor".
not          # Pop one value and push its binary negation (one's complement).

## Notice the distinction between not (flipping all bits) and neg (which
## leaves 0 unchanged). These names are standard in assembly language.

## The pow2 statement pops x (unsigned) and pushes 2 to the power of x.
pow2
shift_l      # Sugar for pow2 mult
shift_ru     # Sugar for pow2 div_u
shift_rs     # Almost sugar for pow2 div_s, except handle the special
              # case of shifting the least negative value 63 bits to the
              # right. (It should be -1.)
```

### ### 9. COMPARISON

```
## For convenience, the comparison predicates below all return 0 for false
## and -1 (all bits set) for true. They also come in sugared variants.

eq           # Pop two values. Push -1 if they are equal, otherwise push 0.
eq! 7        # Pop value and compare it to 7.
```

```

eq!! xx yy    # Sugar for  push!! xx yy eq .

## The remaining comparison operators come in signed and unsigned variants.
lt_u          # Pop x, then y. Push -1 if y < x (unsigned), otherwise push 0.
lt_s          # Pop x, then y. Push -1 if y < x (signed), otherwise push 0.
lte_u         # less than or equal (unsigned)
lte_s         # less than or equal (signed)

gt_u          # y > x (unsigned)
gt_s          # y > x (signed)
gte_u         # greater than or equal (unsigned)
gte_s         # greater than or equal (signed)

### 10. IO STATEMENTS

## To be determined later.

### 11. SET_SP, EXIT

set_sp        # Pop A from the stack, and set the stack pointer to A.
set_sp! 9     # Sugar for  push! 9 set_sp .

exit          # Terminate the machine.

### EXPECTED STACK:
###

```

## Basic assembly

```

### iVM assembly language introduction.
###
### Part 2 - Basics
###
### This is part 2 of the iVM assembly language introduction. Unlike in part 1,
### the statements below should actually make some sense.

### 1. THE STACK

## When the machine starts there are two segments of allocated memory. The
## first segment contains the program. The program counter (PC) initially
## points to the start of this segment. The second segment simply contains
## 24 bytes (3 * 64 bits). The stack pointer (SP) initially points to the
## the first byte after this segment. In other words, we have an initial
## stack size of 3, which is just enough for the program itself to create a
## fresh stack. Since this is something every non-trivial program must do,
## the necessary code is prepended to every executable binary. The default
## stack size is 64 KiB, but it can be changed in the project file (.proj).

## It is often necessary push a copy of an element in the stack onto the
## stack. This can be done as follows:
push!!!! 13 12 11 10    # Push 4 numbers onto the stack (from left to right).
push!! $0 $3           # Push copies of the stack elements 0 and 3 (counting from 0).

## Be careful to use the right number of exclamation marks.
## Otherwise, the parser gets very confused.
## Now the stack is (13, 10, 10, 11, 12, 13) from the top
## since the second statement is sugar for:
push! $0
push! $4                # Notice the offset.

## Stack: (11, 13, 13, 10, 10, 11, 12, 13)

## The previous statement is sugar for:
push! (load8 &4)

## which is sugar for:
push! &4
load8

```

```

## In other words, &n is the address of the element n on the stack.
## This is e.g. useful when we want to pop elements from the stack:
set_sp! &10

## Now the stack is empty again.

### 2. LABELS

## A data segment can be used for shared global variables.
jump! after_x          # Equivalent to jump! (+ x 1)
x:
  data1 [0]
after_x:
  store!! (+ 1 (load1 x)) x # Increase x with 1

## If we export x, it can even be accessed from other files.
EXPORT x

## The store statement above is simply (syntactic) sugar for:
push! (+ 1 (load1 x))
push! x
store1

## Moreover, the first push statement is sugar for:
push! 1
push! (load1 x)
add

## Finally, the second of these statements is sugar for:
push! x
load1

## Now the byte at x is 2, and the stack is (2, 3) from the top.

## Observe that x is not translated into a constant by the assembler.
## Whereas we plan to add an optimization for the initial program (which is
## guaranteed to be located at address 0) the default mode is to produce
## code which is independent of where it is located in memory. In fact, the
## current virtual machine currently randomizes where the code is put in
## order to detect code relying on absolute addresses.

### 3.

## TO BE CONTINUED ...

exit

### EXPECTED STACK:
### 2
### 3

```

## A more advanced example

```

### iVM assembly language introduction.
###
### Part 3 - Advanced
###
### This is part 3 of the iVM assembly language introduction.

### 1. MEMORY MAP WITH NO ENTRY POINT

## In most cases you should specify an entry point when calling the
## assembler, e.g. "-e main". This situation is explained in section 2
## below.

## When an entry point is _not_ specified, execution starts at the first
## statement of the first source file. Moreover, there is code prepended to
## the executable which stores a pointer to the start of the heap in the 8

```



```

    ## bytes preceding the rest of the code:
program:
    heap_start = (load8 (+ program -8))

    ## Similarly, we also have a pointer to the argument file:
    arg_location = (load8 (+ program -16))

    ## More precisely, arg_location is the address of a memory location holding
    ## the length of the argument file. The file contents follows.
    arg_length = (load8 arg_location)
    arg_start = (+ arg_location 8)
    arg_stop = (+ arg_start arg_length)

```

### ### 2. MEMORY MAP WITH ENTRY POINT

```

## If an entry point _was_ specified ("-e main"), then this is where
## execution starts, and the stack looks like this:
##
## 0: return address
## 1: start of argument array
## 2: length of argument array
## 3: start of free heap space

## The entry point is essentially called like this:
push!!! heap_start arg_length arg_start
call! main
set_sp! &2
exit

## The top of the end stack is (truncated and) used as the exit status of
## the virtual machine (unless it is run with "ivm check"). Hence, the
## pointer to the heap start should be overwritten with 0 to signal success.

## Observe that the stack and heap share the same memory. In other words,
## the stack pointer marks the end of the heap.

```

main:

### ### 3. COMPLEX DATA STATEMENTS

```

## Whereas the expressions used to initialize data segments must be
## "assembly time" constants, they may involve both abbreviations and
## labels.
some_constant = 17
jump! after_data
my_data:
    data4 [ (* some_constant (+ after_data -main)) ]
after_data:
    # Set exit status 119 (17*7).
    store8!! (load4 my_data) &3

```

### ### 4. Check binary version compatibility

```

## Ideally, the iVM machine should never change, but it cannot be ruled out.
## If this happens, it is important to know if the machine is compatible
## with the executable binaries at hand. For this reason, it is best
## practice to start assembly files with the following statement:
check_version

## The effect of this statement is to push to the stack the version of the
## machine and instruction set for which the binary was assembled and then
## call the CHECK instruction to see if the machine can run it. If not, the
## machine should signal an error.

```

### ### 5.

```

## To be continued...
return

```

```
### EXPECTED STACK:  
###  
### 119
```